

Battery Health Prediction System - Simple Explanation

Introduction

The objective of this project will experiment with a better, data-driven, way to address this problem. Using machine learning, or Linear Regression model, we have developed a system using the voltage readings of the individual cells of the battery to predict the batteries' state of health. This leads to preventative maintenance and enhanced battery management; saving time, money and potentially reduces risk.

What This System Does

This battery management software estimates the health of a battery pack by monitoring the voltage of each of its individual cells. Imagine it like a doctor checking your vital signs to understand your overall health; here, we are checking the battery's "vital signs" (voltages) to assess the battery's state.

This software works with battery packs which have 21 individual cells. Each cell has a voltage reading (labeled U1 through U21) and the software predicts something called the "State of Health" (SOH) based on all 21 of the voltages. The SOH is a number 0 to 1, where 1.0 represents the battery is new and healthy (100%), while lesser numbers indicate that the battery has degraded over time.

How It Works

The first step in the program is to load a dataset from a CSV file which contains historical battery measurements. The dataset included measurements about 670 different battery packs which contained information about their cell voltages at the time, and what their actual health was. The program then divides that data into two groups, where 80% is used to train the computer model (536 battery packs) and 20% is reserved to test how well the model works (134 battery packs).

The heart of the system uses something called Linear Regression, which is a mathematical technique that identifies patterns in the data. The model is learning the relationship between cell voltage and battery health. After training, it can take a new set of 21 voltage measurements and predict that new battery packs' state of health (SOH) readings, even if it never sees that specific battery pack before.

Understanding the Results

When you executed the program, it indicated the model achieved an R^2 score of 0.6561, meaning it accounted for roughly 66% of what governs battery health. This is okay; it's not fantastic, but it's sufficiently good to be useful. Specifically, the Mean Absolute Error of 0.0303 indicates the predictions are off by about 3% on average, which is relatively accurate for practical purposes in battery health assessments.

The visualization graph you reviewed lays the model's predictions (red line) against the actual known values (blue line). They resemble a similar pattern; however, the red line does not correlate completely to the blue line. In sum, this demonstrates the model is somewhat functional, but clearly, there is room for improvement. Finally, the program demonstrated a real-world example by taking average voltage readings, which predicted a battery pack with those average readings would have a health of about 81.67% of its original health.


Real-World Use

Such a system would be invaluable in electric vehicles, smartphones, or solar energy storage systems. Rather than waiting for the battery to fully fail, the system forecasts that the battery is degrading and alerts you in advance. This supports planning for maintenance, lets you avoid failure, and alerts you getting replacement batteries when you should in anticipation of failure.

Conclusion

This battery health prediction system exemplifies how machine learning can be used practically for real-world engineering problems. With about 97% accuracy (or 3% average error), the model can offer reliable forecasting tools that can be relied on for any critical maintenance decision making process. The code structure is neat and modular allowing for usability, maintainability, and enhancement possibilities. The system can certainly be enhanced, perhaps a more complex algorithm. Or adding features such as temperature or charge cycles, but it already establishes value. It takes voltage readings, moves through the model to successively develop insight, aiding in management decisions towards what is affordable and minimum safety practices for longer untreated battery use.

OUTPUT:



The screenshot shows a Jupyter Notebook interface with two main sections. On the left, a figure titled "Actual vs Predicted SOH Across Test Samples" displays a scatter plot with a blue line for "Actual SOH" and a red line for "Predicted SOH". The x-axis is "Test Sample Index (sorted by actual SOH)" ranging from 0 to 120, and the y-axis is "SOH Value" ranging from 0.65 to 0.95. On the right, the corresponding Python code is shown:

```
mean_absolute_error, mean_squared_error

# Load dataset
df = pd.read_csv("PulseBat Dataset.csv")

# Split dataset into features and target
X = df.drop("SOH", axis=1)
y = df["SOH"]

# Split into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)


# Make predictions
y_pred = model.predict(X_test)

# Evaluate model performance
mean_absolute_error(y_test, y_pred), mean_squared_error(y_test, y_pred)
```

Below the code, the terminal output shows the results of the linear regression model training and evaluation:

```
Testing set: 134 samples
Training Linear Regression model...
=====
MODEL PERFORMANCE RESULTS
=====
R² Score: 0.6561
Mean Absolute Error (MAE): 0.0303
Root Mean Squared Error (RMSE): 0.0387
```

At the bottom, the system tray shows the date and time as 2025-10-23 8:50 PM.



The screenshot shows a Jupyter Notebook interface with the following details:

- File Explorer:** Shows the project structure under "BATTERY-MANAGEMENT-SYSTEM": data (PulseBat Dataset.csv, PulseBat Dataset.xlsx) and src (LinearRegression.py).
- Terminal:** Displays the command: PS C:\Users\mdbur\Desktop\battery-management-system\src> python LinearRegression.py. The output shows the loading of data, training of the model, and the resulting model performance metrics (R² Score: 0.6561, MAE: 0.0303, RMSE: 0.0387).
- Output:** Shows the execution of the script and the resulting model performance results.
- Bottom Status Bar:** Shows the date and time as 2025-10-23 8:50 PM.

The screenshot shows a dark-themed IDE interface with a terminal window open. The terminal window displays the output of a Python script named `LinearRegression.py`. The script performs a battery management analysis, specifically a linear regression analysis on battery voltage data. Key outputs include:

- U21: 3.708V
- Predicted SOH for this battery pack: 0.8167
- ANALYSIS COMPLETE
- U18: 3.770V
U19: 3.817V
U20: 3.747V
U21: 3.708V
- Predicted SOH for this battery pack: 0.8167
- ANALYSIS COMPLETE
- U18: 3.770V
U19: 3.817V
U20: 3.747V
U21: 3.708V
- Predicted SOH for this battery pack: 0.8167
- ANALYSIS COMPLETE
- U18: 3.770V
U19: 3.817V
U20: 3.747V
U21: 3.708V
- Predicted SOH for this battery pack: 0.8167
- ANALYSIS COMPLETE
- U18: 3.770V
U19: 3.817V

The terminal window also shows the command prompt `PS C:\Users\mdbur\Desktop\battery-management-system\src>`. The status bar at the bottom indicates the script is running in Python 3.14.0, with 8:51 PM and 2025-10-23.

This screenshot shows the same IDE environment after the script has completed its execution. The terminal window now displays the final output of the analysis, which includes the predicted SOH value of 0.8167 and the analysis complete message. The command prompt is shown as `PS C:\Users\mdbur\Desktop\battery-management-system\src>`. The status bar at the bottom indicates the script is running in Python 3.14.0, with 8:51 PM and 2025-10-23.